

Introduction to MapReduce

Today's Topics

- Functional programming
- MapReduce
- Distributed file system

Functional Programming

- MapReduce = functional programming meets distributed processing on steroids
 - Not a new idea... dates back to the 50's (or even 30's)
- What is functional programming?
 - Computation as application of functions
 - Theoretical foundation provided by lambda calculus
- How is it different?
 - Traditional notions of “data” and “instructions” are not applicable
 - Data flows are implicit in program
 - Different orders of execution are possible
- Exemplified by LISP and ML

Overview of Lisp

- Lisp ≠ Lost In Silly Parentheses
- We'll focus on particular a dialect: "Scheme"
- Lists are primitive data types
 - `'(1 2 3 4 5)`
 - `'((a 1) (b 2) (c 3))`
- Functions written in prefix notation
 - `(+_1 2) → 3`
 - `(* 3 4) → 12`
 - `(sqrt (+ (* 3 3) (* 4 4))) → 5`
 - `(define x 3) → x`
 - `(* x 5) → 15`

Functions

- Functions = lambda expressions bound to variables
- ```
(define foo
 (lambda (x y)
 — (sqrt (+ (* x x) (* y y)))))
 —
```
- Syntactic sugar for defining functions
  - Above expressions is equivalent to:  

```
(define (foo x y)
 (sqrt (+ (* x x) (* y y))))
—
```
- Once defined, function can be applied:
- ```
(foo 3 4) → 5
```

Other Features

- In Scheme, everything is an s-expression
 - No distinction between “data” and “code”
 - Easy to write self-modifying code
- Higher-order functions
 - Functions that take other functions as arguments

```
(define (bar f x) (f (f x)))
```

Doesn't matter what *f* is, just apply it twice.

```
(define (baz x) (* x x))
```

```
(bar baz 2) → 16
```

Recursion is your friend

- Simple factorial example

- ```
(define (factorial n)
 (if (= n 1)
 1
 (* n (factorial (- n 1)))))
```
- ```
(factorial 6) → 720
```

- Even iteration is written with recursive calls!


```
(define (factorial-iter n)
  (define (aux n top product)
    (if (= n top)
        (* n product)
        (aux (+ n 1) top (* n product))))
  (aux 1 n 1))
(factorial-iter 6) → 720
```

Lisp → MapReduce?

- What does this have to do with MapReduce?
- After all, Lisp is about processing *lists*
- Two important concepts in functional programming
 - Map: do something to everything in a list
 - Fold: combine results of a list in some way


Map

- Map is a higher-order function
- How map works:
 - Function is applied to every element in a list
 - Result is a new list



Fold

- Fold is also a higher-order function
- How fold works:
 - Accumulator set to initial value
 - Function applied to list element and the accumulator
 - Result stored in the accumulator
 - Repeated for every item in the list
 - Result is the final value in the accumulator



Map/Fold in Action

- Simple map example:

- `(map (lambda (x) (* x x))`
- `'(1 2 3 4 5))`
- `→ '(1 4 9 16 25)`
-

- Fold examples:

- `(fold + 0 '(1 2 3 4 5)) → 15`
- `(fold * 1 '(1 2 3 4 5)) → 120`

- Sum of squares:

```
(define (sum-of-squares v)
  (fold + 0 (map (lambda (x) (* x x)) v)))
(sum-of-squares '(1 2 3 4 5)) → 55
```

Lisp → MapReduce

- Let's assume a long list of records: imagine if...
 - We can parallelize map operations
 - We have a mechanism for bringing map results back together in the fold operation
- That's MapReduce! (and Hadoop)
- Observations:
 - No limit to map parallelization since maps are independent
 - We can reorder folding if the fold function is commutative and associative

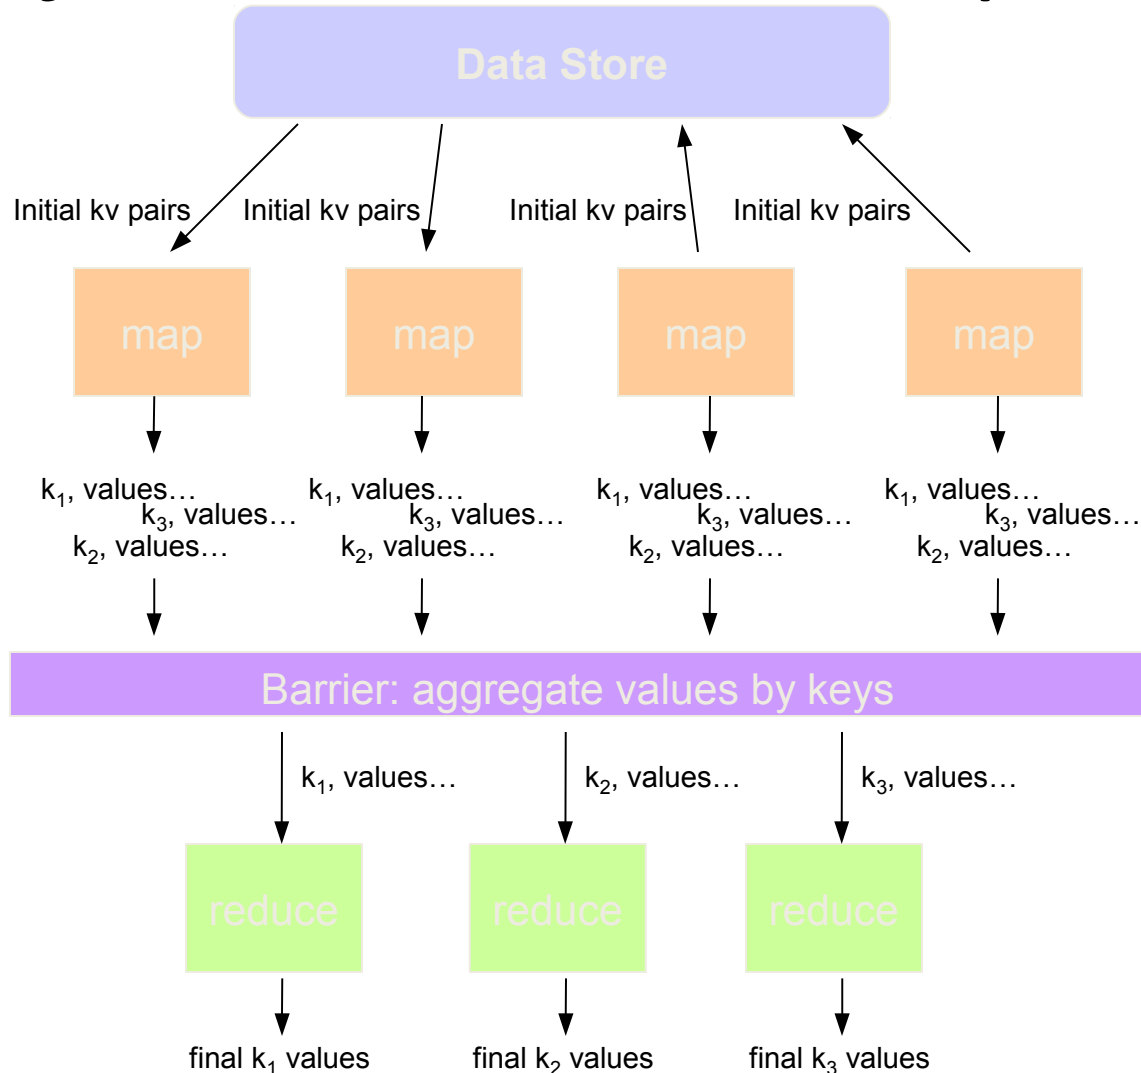
Typical Problem

- Iterate over a large number of records
- Map* • Extract something of interest from each
- Shuffle and sort intermediate results
- Aggregate intermediate results *Reduce*
- Generate final output

MapReduce

- Programmers specify two functions:
 - map** $(k, v) \rightarrow \langle k', v' \rangle^*$
 - reduce** $(k', v') \rightarrow \langle k', v' \rangle^*$
 - All v' with the same k' are reduced together
- Usually, programmers also specify:
 - partition** $(k', \text{number of partitions}) \rightarrow \text{partition for } k'$
 - Often a simple hash of the key, e.g. $\text{hash}(k') \bmod n$
 - Allows reduce operations for different keys in parallel
- Implementations:
 - Google has a proprietary implementation in C++
 - Hadoop is an open source implementation in Java (lead by Yahoo)

It's just divide and conquer!



Recall these problems?

- How do we assign work units to workers?
- What if we have more work units than workers?
- What if workers need to share partial results?
- How do we aggregate partial results?
- How do we know all the workers have finished?
- What if workers die?

