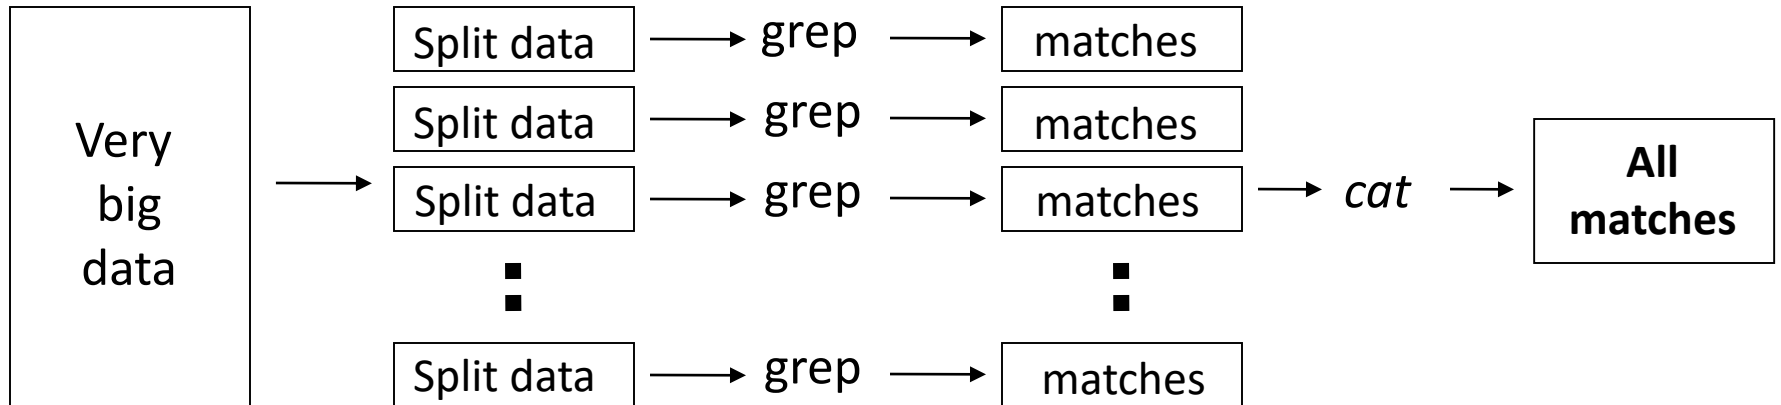


Map Reduce

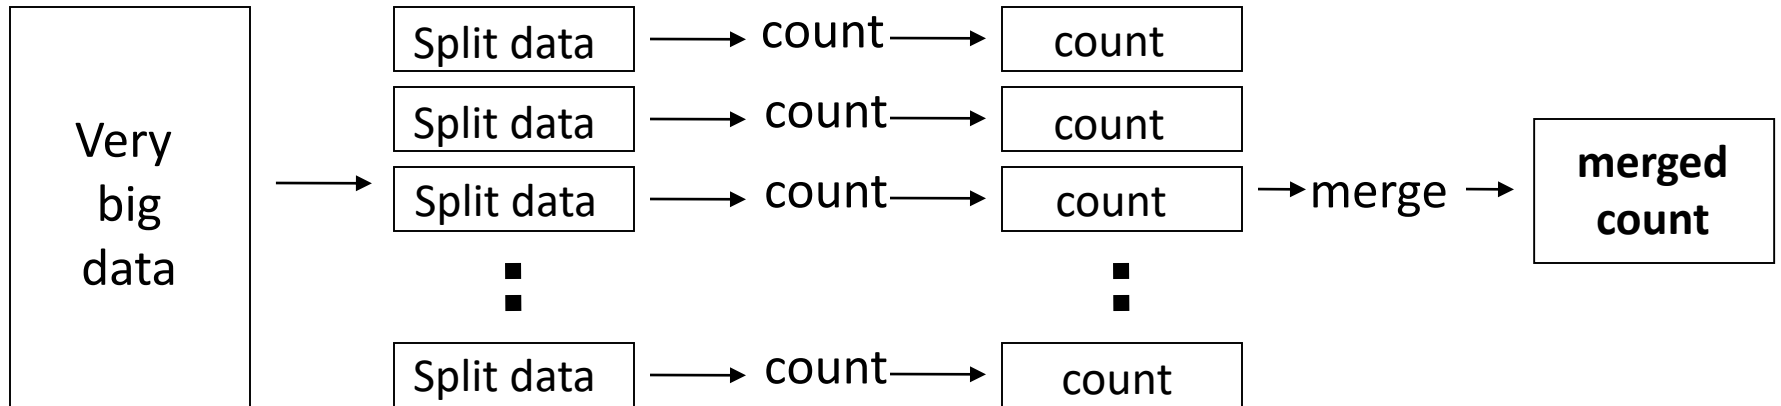
What is MapReduce?

- A programming model (& its associated implementation)
- For processing large data set
- Exploits large set of commodity computers
- Executes process in distributed manner
- Offers high degree of transparencies

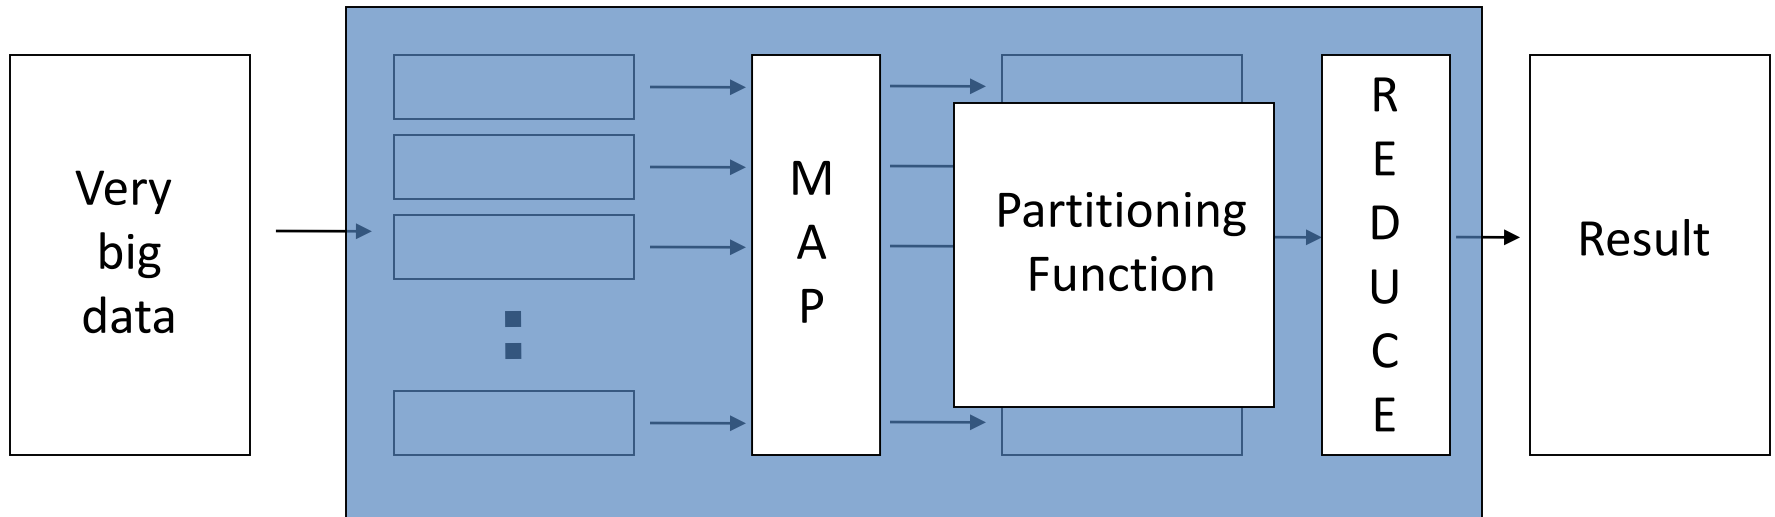
Distributed Grep



Distributed Word Count



Map Reduce



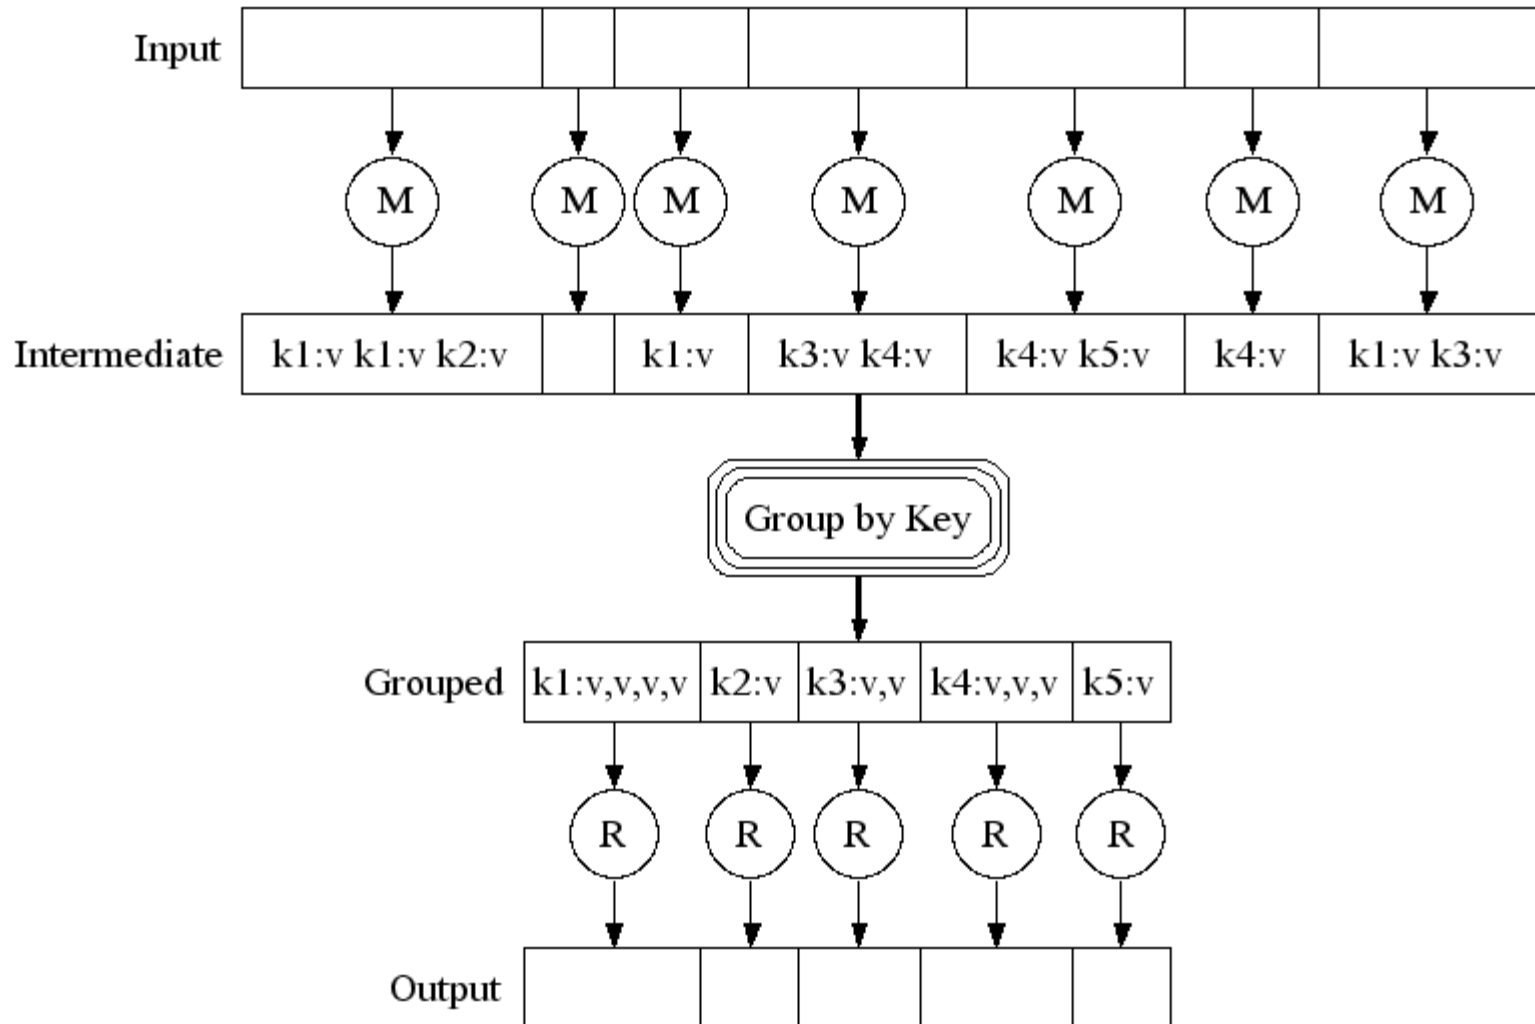
- Map:

- Accepts *input* key/value pair
- Emits *intermediate* key/value pair

- Reduce :

- Accepts *intermediate* key/value* pair
- Emits *output* key/value pair

Partitioning Function



Example for MapReduce

- Page 1: the weather is good
- Page 2: today is good
- Page 3: good weather is good.

Map output

- Worker 1:
 - (the 1), (weather 1), (is 1), (good 1).
- Worker 2:
 - (today 1), (is 1), (good 1).
- Worker 3:
 - (good 1), (weather 1), (is 1), (good 1).

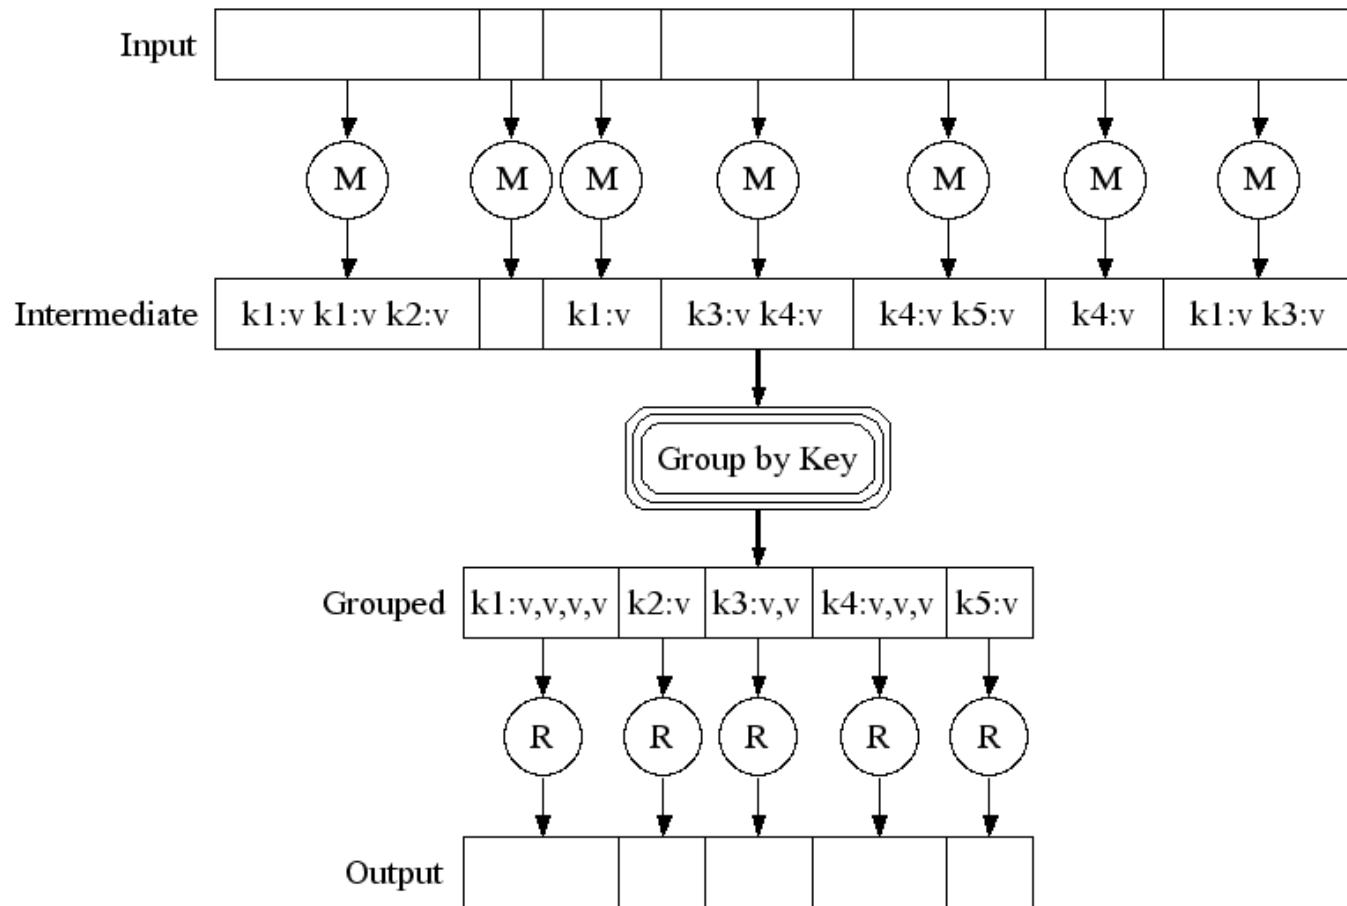
Reduce Input

- Worker 1:
 - (the 1)
- Worker 2:
 - (is 1), (is 1), (is 1)
- Worker 3:
 - (weather 1), (weather 1)
- Worker 4:
 - (today 1)
- Worker 5:
 - (good 1), (good 1), (good 1), (good 1)

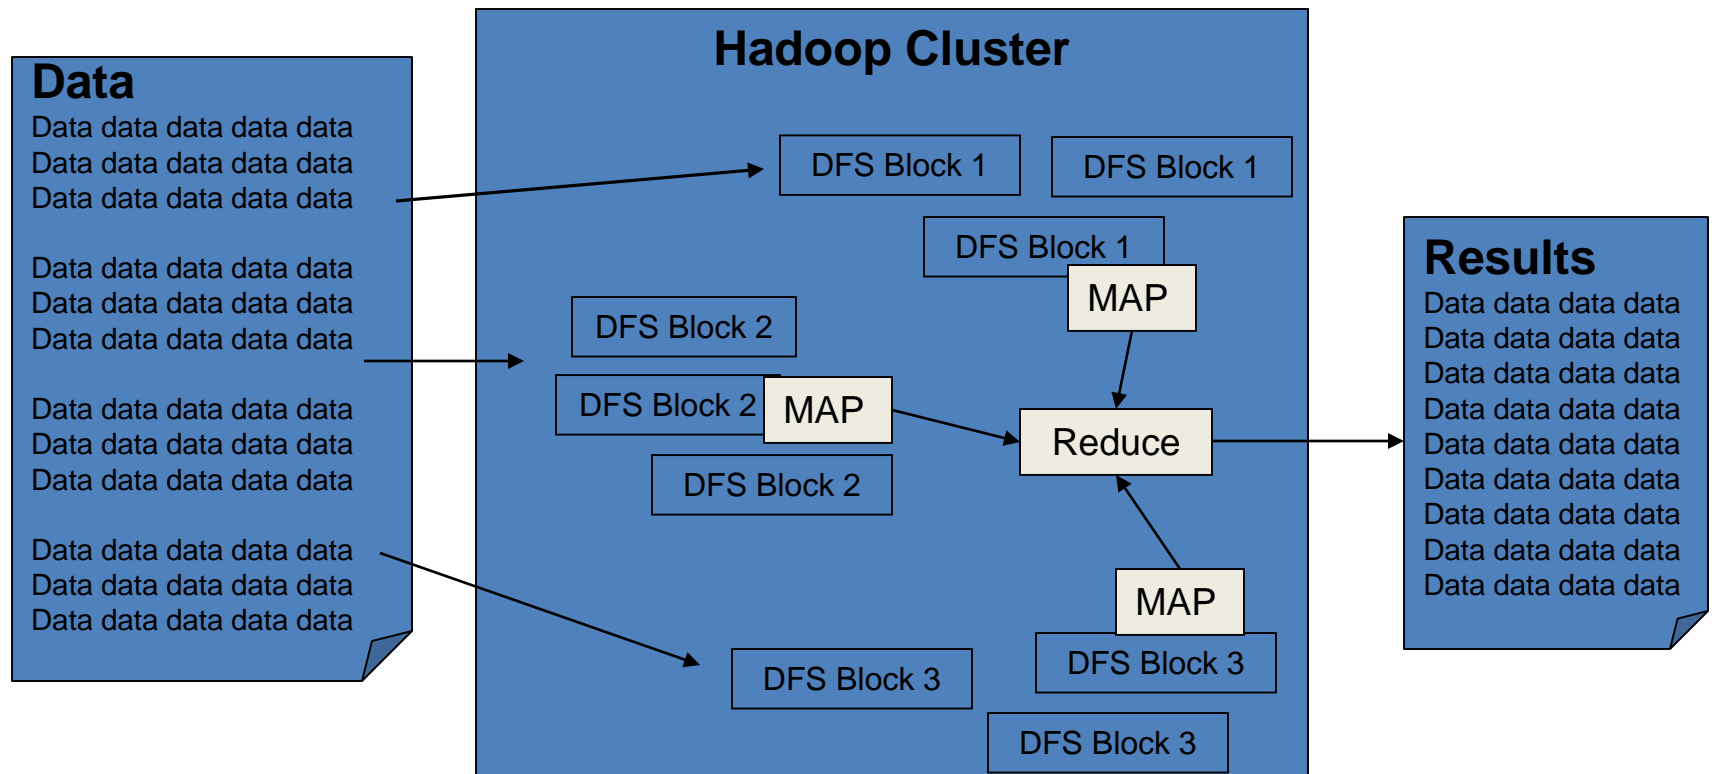
Reduce Output

- Worker 1:
 - (the 1)
- Worker 2:
 - (is 3)
- Worker 3:
 - (weather 2)
- Worker 4:
 - (today 1)
- Worker 5:
 - (good 4)

MapReduce Architecture



Hadoop Architecture



Sample Hadoop Code

- **Sample text-files as input:**

- ```
$ bin/hadoop dfs -ls /usr/joe/wordcount/input/
/usr/joe/wordcount/input/file01
/usr/joe/wordcount/input/file02
```

```
$ bin/hadoop dfs -cat /usr/joe/wordcount/input/file01
Hello World, Bye World!
```

```
$ bin/hadoop dfs -cat /usr/joe/wordcount/input/file02
Hello Hadoop, Goodbye to hadoop.
```

- **Run the application:**

- ```
$ bin/hadoop jar /usr/joe/wordcount.jar org.myorg.WordCount  
/usr/joe/wordcount/input /usr/joe/wordcount/output
```

- **Output:**

- ```
$ bin/hadoop dfs -cat /usr/joe/wordcount/output/part-00000
Bye 1
Goodbye 1
Hadoop, 1
Hello 2
World! 1
World, 1
hadoop. 1
to 1
```

# Contd...

- Notice that the inputs differ from the first version we looked at, and how they affect the outputs.
- Now, lets plug-in a pattern-file which lists the word-patterns to be ignored,
- ```
$ hadoop dfs -cat /user/joe/wordcount/patterns.txt
```

```
\.  
\!  
to
```
- Run it again, this time with more options:
- ```
$ bin/hadoop jar /usr/joe/wordcount.jar org.myorg.WordCount –
```
- ```
Dwordcount.case.sensitive=true /usr/joe/wordcount/input
```
- ```
/usr/joe/wordcount/output -skip /user/joe/wordcount/patterns.txt
```
- As expected, the output:
- ```
$ bin/hadoop dfs -cat /usr/joe/wordcount/output/part-00000
```

```
Bye 1  
Goodbye 1  
Hadoop 1  
Hello 2  
World 2  
hadoop 1
```

Contd...

- Run it once more, this time switch-off case-sensitivity:
- `$ bin/hadoop jar /usr/joe/wordcount.jar
org.myorg.WordCount -Dwordcount.case.sensitive=false
/usr/joe/wordcount/input /usr/joe/wordcount/output -skip
/usr/joe/wordcount/patterns.txt`
- Sure enough, the output:
- `$ bin/hadoop dfs -cat /usr/joe/wordcount/output/part-00000`
bye 1
goodbye 1
hadoop 2
hello 2
world 2